Improving Structured Data ML Process with GANs

3/30/21 11:10 AM

Speaker: Srinivas Chilukuri (ZS Associates)
Date: 03/30/2021

Structured datasets have unique challenges of their own

- Capture biases
- Missing data
- Small data
- Class imbalance

GAN overview for image generation

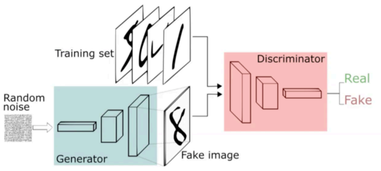

Applications of GANs

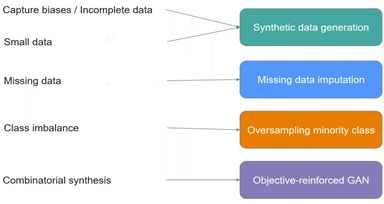

Synthetic generation of tabular data

- Generator create values according to schema of columns
- Discriminator produces log probability

Domain specific tabular GANs

- E-commerce:
  - Generate low-dimensional dense representation of e-commerce orders
  - To generate plausible orders data for new products
- Insurance: (tackle data privacy issue)
  - Generate accessible insurance datasets for actuarial studies
  - Tweaked CTGAN for imbalance categorical levels
- Medical (medGAN)
  - Generate privacy preserving, accessible healthcare datasets
  - Use of recurrent Encoders/Decoders for handling sequences
  - Does not handle lab values, clinical notes, diagnostic images, …
- Relational datasets

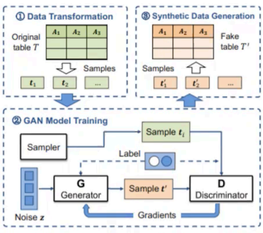

Figure 2: Overview of data synthesis using GAN. (1) It transforms each record in a relational table into a sample $t \in \mathbb{R}^d$. (2) It takes the samples as input to train a deep generative model G using the adversarial training framework in GAN. (3) It utilizes the trained G to generate a set of synthetic samples, which are then transformed back into fake records.

Missing data imputation

- Through masking data points in original dataset

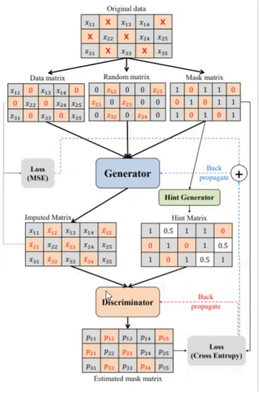

Generate data based on an objective

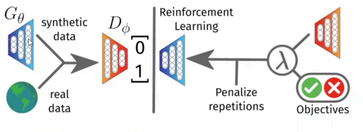

Figure 1: Schema for ORGAN. *Left*: D is trained as a classifier receiving as input a mix of real data and generated data by G. *Right*: G is trained by RL where the reward is a combination of D and the objectives, and is passed back to the policy function via Monte Carlo sampling. We penalize non-unique sequences.

Future directions

1. *How to evaluate tabular GANs?*

   Sample similarity
   - Basic statistics
   - Column correlations
   - Mirror column associations
   - PCA variance correlation

   Machine learning efficacy
   - Evaluate how downstream ML model performance varies with real data, synthetic data, real data + synthetic

   Privacy evaluations
   - Check how many samples are replicated from the training (real) dataset

   Human evaluations
   - On the lines of Turing test - to see whether humans can distinguish between real and synthetic data
   - In domain specific implementations, this will help with understanding which patterns GANs are able to learn

2. *Open challenges*
   - Data hungry - Adaptive Discriminator Augmentation
   - Mode collapse for imbalanced categorical data
   - Domain specific data still requires significant work e.g. HER data is still not realistic